A Regime Theory of Joy

The Ease Regime as a Permissive Control Configuration

Florian Morin
Independent Researcher in Affective Science
florianmorin.com
ORCID: 0009-0004-7301-8524

April 17, 2026

Abstract

Joy is not well captured by accounts that treat it as a direct product of reward processes, learning, or goal-directed activity. Instead, it depends on access conditions and is gated by evaluative monitoring, becoming available only when such monitoring falls below a critical threshold.

This transition defines a distinct control regime, here termed the ease regime. Access to this regime is governed by a state variable, Z, which indexes the level of evaluative monitoring and optimization. Although presented as a single variable, Z reflects the combined influence of multiple interacting, and not necessarily consciously accessible, sources of evaluative load. When Z is high, cognition operates in an evaluation-dominated mode, characterized by continuous correction, goal orientation, and performance tracking, processes that actively constrain affective intensity.

When Z falls below a critical threshold, evaluative processes lose their stabilizing influence over cognition. Discrepancies remain informational rather than being converted into corrective signals, and high-intensity positive affect becomes accessible.

Entry into this regime is abrupt and unambiguous, rather than the result of gradual improvement, supporting the interpretation of a regime shift rather than a learning process. A central blocking mechanism, Z_{shift}, can be understood as a self-reinforcing dynamic within this system, in which attempts to produce the state induce persistent monitoring, thereby preventing entry. The model further predicts a measurement paradox: methods that recruit self-monitoring tend to suppress the very phenomenon they attempt to observe.

Overall, the framework provides a falsifiable alternative to reward-based and training-based accounts by formalizing joy as a threshold-dependent regime of system dynamics, rather than as an outcome of optimization.

Key points

- High-intensity positive affect depends on access to a control regime defined by reduced evaluative coupling, rather than reward or optimization processes.
- Entry is threshold-dependent and is disrupted by monitoring, evaluation, and goal-directed control.
- Prediction error contributes directly to experiential intensity only within this regime.
- The regime enables, but does not require, high-fidelity reactivation of previously encoded experiences.
- Memory reactivation is a consequence of regime access, not its primary mechanism.
- High-intensity positive affect requires both reduced evaluative coupling and sufficient coherence of discrepancy signals to support their propagation.

1 Introduction

Contemporary models of positive affect predominantly conceptualize joy as the outcome of reward processing, goal attainment, or affective valuation (Kent Berridge & Morten Kringelbach, 2015; Wolfram Schultz, 2016). Within this framework, increases in positive affect are typically treated as continuous, scalable, and directly producible through identifiable inputs such as reinforcement, prediction error, or cognitive appraisal (Richard Sutton & Andrew Barto, 2018; Robb Rutledge et al., 2014; Daniel Kahneman, 1999; Ed Diener, 2000). Parallel accounts in appraisal theory similarly treat positive affect as the output of evaluative interpretations of events relative to goals and expectations (Klaus Scherer, 2009; Magda Arnold, 1960). While contemporary models increasingly refine the computational structure of reward processing, including recent accounts of dopaminergic prediction error and subjective well-being, these approaches continue to treat positive affect as a graded and directly producible output of identifiable variables (Berridge & Kringelbach, 2015; Schultz, 2016; Rutledge et al., 2014; see also recent extensions in Berridge, 2018; Schultz, 2021).

However, this view struggles to account for a class of phenomenological observations in which high-intensity positive states appear abruptly, resist voluntary induction, and collapse under attempts at control or stabilization, as reported in initial observational work (Morin, 2026). Related anomalies have been noted in studies of flow, awe, and peak experience, where transitions are often described as discrete and difficult to volitionally reproduce (Mihaly Csikszentmihalyi, 1990; Abraham Maslow, 1964; Dacher Keltner & Jonathan Haidt, 2003). These observations suggest that certain forms of positive affect are not well described as outputs of additive reward mechanisms, but instead reflect access to a distinct regime of cognitive-affective organization. In such cases, the transition into the state is discrete rather than gradual, often described as a sudden shift in the structure of experience, accompanied by a qualitative reorganization in the processing of discrepancies, goals, and sensory inputs.

Notably, these states cannot be reliably produced through effort, training, or optimization strategies, and are often disrupted by the very processes typically assumed to generate them. This paradox resonates with findings in cognitive science showing that introspection and control can interfere with ongoing experience, particularly in domains requiring fluid, non-evaluative processing (Jonathan Schooler, 2002; Timothy Wilson & Jonathan Schooler, 1991; Sian Beilock, 2010). Similar disruptions have been documented in motor expertise and creativity, where explicit monitoring degrades performance (Anders Ericsson, 2006; Mark Beeman & John Kounios, 2009). More recent work on meta-consciousness and introspective access further supports this constraint, suggesting that attempts to explicitly monitor ongoing experience can alter or degrade its underlying structure (Schooler, 2002; see also Schooler, 2011; Epley & Whitchurch, 2008).

A central difficulty in studying such states lies in a methodological constraint that has remained largely under-addressed. Standard approaches to affect measurement, including self-report scales and task-based evaluation, implicitly introduce evaluative demands that alter the underlying cognitive regime. These procedures require subjects to monitor, interpret, and report their internal states, thereby engaging control processes that may prevent the emergence of the very phenomena under investigation. This aligns with broader concerns about reactivity and observer effects in psychological measurement (Donald Campbell, 1957; Harold Kelley, 1967).

As a result, the absence or rarity of such states in experimental settings may not reflect their true accessibility, but rather the conditions imposed by the measurement process itself. This leads to what can be termed a measurement paradox in affect research: the more precisely a high-intensity positive state is targeted or operationalized, the lower its probability of occurrence. Related tensions have been noted in neuroscience, where attempts to stabilize or repeatedly induce affective responses often lead to habituation or attenuation (Kent Berridge, 2003; Wolfram Schultz, 2016), and in predictive processing accounts where increasing model precision can suppress exploratory or affectively rich states (Karl Friston, 2010; Andy Clark, 2013).

This issue is not limited to extreme or rare experiences, but may extend to common developmental phenomena, such as the apparent accessibility of intense positive affect during childhood and its progressive attenuation under adult cognitive constraints (Alison Gopnik, 2009; Paul Harris, 2000).

In response to these limitations, the present paper proposes a regime-based account of joy. Rather than treating positive affect as a scalar output of reward systems, joy is conceptualized as emerging from a specific configuration of cognitive control in which discrepancy signals remain informational and are not immediately converted into corrective demands. This configuration defines a permissive regime, characterized by reduced evaluative stabilization and altered dynamics of prediction error processing

(Friston, 2010; Clark, 2013). Within this view, entry into joy is governed by threshold dynamics rather than accumulation. The transition occurs when evaluative control fails to stabilize discrepancies, allowing a shift in the functional role of prediction error from corrective signal to experiential driver.

This model generates distinct empirical predictions, including abrupt onset, all-or-none phenomenology, sensitivity to contextual evaluation load, and incompatibility with sustained goal-directed optimization. By reframing joy as a regime shift rather than a reward output, this account aims to resolve inconsistencies in the existing literature and to provide a basis for new experimental approaches that do not interfere with the phenomena they seek to measure.

2 Reframing "Ease" as a Computational Regime

The term "ease" is used here as a provisional label for a specific configuration of cognitive-affective dynamics. This account does not introduce a new primitive category of affect, but identifies a configuration of control parameters that gives rise to a distinct class of experiential states. This distinctiveness is emergent rather than fundamental. A regime is defined as a stable configuration of relationships between key functional components, including discrepancy detection, evaluative processes, and action selection. Regimes differ not by the presence or absence of these components, but by how they are coupled. The present account focuses on a configuration in which discrepancy signals are not immediately converted into corrective or goal-directed processes. Under standard conditions, discrepancies between expected and observed states trigger rapid evaluation and correction, supporting efficient goal pursuit but imposing a continuous evaluative structure on experience. In contrast, under reduced evaluative coupling, discrepancies persist as informational signals rather than initiating control processes.

This shift has three main consequences. First, discrepancies are no longer rapidly resolved, altering the temporal dynamics of processing. Second, prediction error contributes directly to experiential intensity rather than functioning solely as a transient control signal. Third, the dominance of goal-directed structure is reduced, allowing perception and action to unfold without continuous optimization pressure. Importantly, this regime does not reflect a deficit or pathological failure of control. The underlying architecture remains intact, and evaluative processes can be re-engaged at any moment. Rather, the regime reflects a change in control policy, specifically a reduction in the gain assigned to evaluative processes. Accordingly, "ease" does not denote a separate affective category, but a configuration within a broader space of possible control regimes. Access to the permissive regime depends on reduced evaluative coupling, while the emergence of high-intensity positive affect additionally requires that discrepancy signals remain sufficiently coherent to propagate within the system.

The present proposal isolates one such configuration without claiming phenomenological exclusivity. Its contribution lies in situating the phenomenon within established computational frameworks, including predictive processing, control theory, and reinforcement learning, while allowing it to be described in terms of parameter changes and coupling dynamics rather than as an exceptional category. The regime is not defined by specific content, including memory. Its defining property is reduced evaluative coupling. Under these conditions, two classes of phenomena may emerge. First, high-intensity positive affect becomes accessible as a regime property. Second, previously encoded experiential configurations may be reactivated with high fidelity. Such reactivation is contingent on regime access and should be interpreted as a consequence of the underlying control dynamics rather than as a primary mechanism. Z is a composite, macro-level index of evaluative load, while g denotes the local coupling parameter through which this load is expressed in the transformation of discrepancy into corrective signals.

2.1 Z definition

We distinguish three levels of description.

- Monitoring refers to a control-layer process that continuously tracks discrepancies between expected and observed states.

- Evaluation and optimization are behavioral consequences of sustained monitoring, reflecting the systematic conversion of discrepancies into corrective, goal-directed adjustments.

- Z is defined as a macro-level index of this evaluative load, integrating multiple sources of monitoring across timescales. This load is implemented locally through a coupling parameter g, which governs the transformation of discrepancy signals into corrective drives.

Accordingly, the relationship between these components can be summarized as follows: monitoring → evaluative load (Z) → coupling strength (g) → conversion of discrepancy into correction.

3 Relation to Existing Frameworks

The regime proposed here overlaps superficially with several well-established constructs in the study of attention, affect, and consciousness. However, these similarities are primarily phenomenological and do not extend to the underlying dynamics or access conditions. A comparative analysis helps clarify these distinctions. All phenomena described in this framework are derived from a single control relationship: evaluative monitoring increases evaluative load (Z), which modulates coupling strength (g), thereby determining whether discrepancy signals are rapidly corrected or allowed to persist.

3.1 Meditative absorption (e.g., jhāna states

Meditative absorption states are typically characterized by sustained attention, progressive stabilization, and the cultivation of specific attentional skills over extended training periods. Entry into such states is often gradual and depends on the active maintenance of attentional focus. By contrast, the regime described here does not require training or sustained attentional control. Entry is abrupt and cannot be achieved through deliberate stabilization. In fact, attempts to maintain or reproduce the state through controlled attention tend to disrupt it. This suggests that, while both involve altered experiential structure, they rely on different control policies: one based on attentional stabilization, the other on a reduction of evaluative monitoring.

3.2 Flow states

Flow is commonly defined as a state of optimal engagement during goal-directed activity, where skill and challenge are well matched. It involves high task focus, efficient performance, and a tight coupling between perception and action within a structured objective. The present regime differs in that it is not tied to task optimization or performance. It can occur in the absence of goals and is often disrupted by the introduction of explicit objectives. Rather than reflecting efficient control within a task, it reflects a temporary suspension of optimization pressures altogether. As such, it is anti-instrumental in its access conditions, even if it may coexist with ongoing activity

4 Reward-based and Reinforcement Learning Accounts

Standard accounts of positive affect link pleasure to reward prediction error, reinforcement signals, or dopaminergic valuation processes. Within these models, affect is typically treated as a function of outcomes, expectations, and their discrepancies.

In the present framework, prediction error is not eliminated but persists under reduced evaluative coupling, contributing directly to experiential intensity rather than being rapidly converted into corrective signals. The key difference lies not in its presence, but in its mode of integration within the control architecture.

This distinction can be decomposed into two roles. Prediction error provides the energetic component of the experience, while reduced evaluative coupling defines the access condition. Accordingly, prediction error is necessary but not sufficient for the emergence of high-intensity positive affect, as its impact depends on the underlying control regime. Consistent with this view, dopaminergic signals are widely interpreted as encoding prediction error rather than subjective pleasure per se (Schultz, 2016, 2021), reinforcing the separation between computational learning signals and phenomenological experience.

Prediction error should therefore be understood as a carrier of experiential intensity within the permissive regime, rather than as its defining mechanism, which is determined by the control architecture mediated by evaluative load (Z) and coupling (g).

These considerations suggest that prediction error alone cannot account for high-intensity affective states, but instead undergoes a control-dependent transformation within the proposed framework.

5 From Prediction Error to Affective Experience in the Permissive Regime

These effects are modulatory rather than generative, as they operate only once the control regime has shifted. In more recent formulations, increased model precision and control have been associated with reduced exploratory dynamics, suggesting that heightened evaluative constraint may suppress modes of experience characterized by open-ended or high-variability signal propagation (Clark, 2019). Within the permissive regime, certain classes of sensory input appear to modulate the intensity of positive affect without being sufficient to induce the state. In particular, visual features associated with high salience, such as brightness, contrast, or properties resembling rarity cues (e.g., glow-like or "treasure-like" signals), are frequently reported to amplify subjective intensity. These features may act by increasing the persistence or propagation of discrepancy signals under reduced evaluative coupling. A second, partially distinct mode of positive affect is also observed. This mode is typically associated with slow-tempo, low-frequency auditory input, often involving simple or minimally structured melodic patterns. It is described as less localized and less sharply defined than the primary mode, and may involve diffuse, low-contrast perceptual qualities. Importantly, these inputs do not reliably produce the regime in isolation. Their effects appear contingent on prior access to the regime and should therefore be interpreted as modulatory rather than generative. This distinction suggests that different input classes may interact with the regime through partially dissociable pathways, potentially corresponding to different modes of discrepancy propagation.

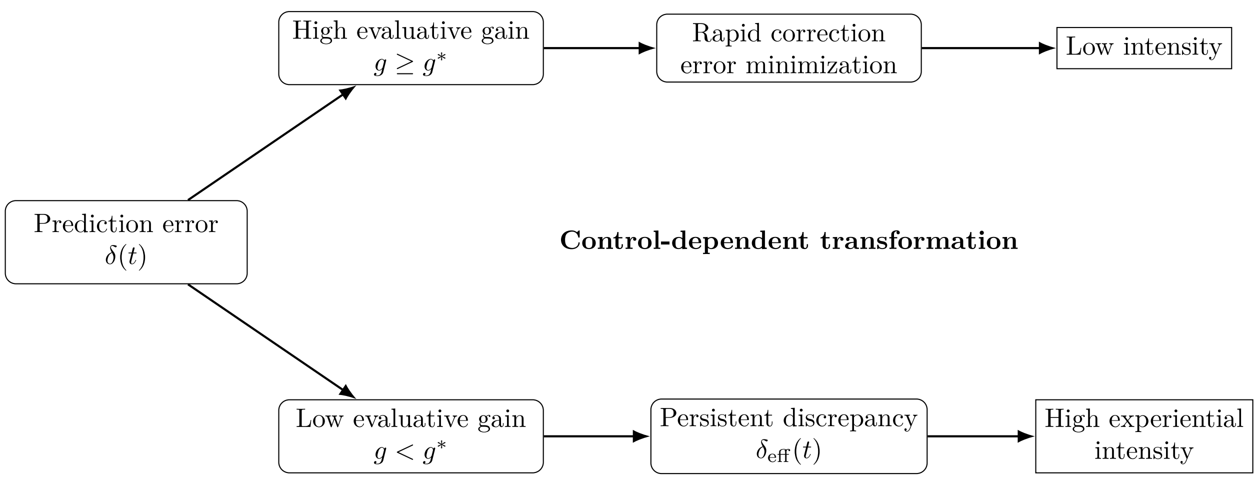

Figure 1: Control-dependent transformation of prediction error. The same discrepancy signal (δ) leads to different outcomes depending on evaluative gain (g). Under high gain, discrepancies are rapidly corrected and have limited experiential impact. Under low gain, discrepancies persist and accumulate, contributing directly to subjective intensity. This dissociation explains why prediction error alone is insufficient to generate high-intensity positive affect.

6 This is Not Nostalgia: Reactivation as a Downstream Consequence of Regime Access

The following section describes a downstream consequence of regime access, not the core mechanism of the state itself. Within the present framework, the permissive regime does not originate from memory processes, but may enable the reactivation of previously encoded experiential configurations.

Under these conditions, reactivation can involve configurations associated with early-life positive affect. This process is not equivalent to standard autobiographical recall, but reflects a reinstantiation of the original affective and perceptual structure of the experience. For example, exposure to a previously encountered object (e.g., a childhood-related stimulus) may be accompanied by a rapid and coherent re-emergence of its associated affective tone. This reactivation appears to involve the joint reinstatement of sensory, affective, and bodily components, rather than the retrieval of abstract or narrative memory content.

Importantly, this process does not depend on voluntary recall and cannot be reliably induced through deliberate effort. In some cases, it may occur automatically and resist suppression, consistent with a system-level change in access conditions rather than a controlled retrieval process.

From a mechanistic perspective, this phenomenon is consistent with processes resembling hippocampal pattern completion, in which partial cues trigger the reactivation of previously encoded configurations. Under standard evaluative conditions, such reactivation may be attenuated or fragmented by rapid categorization and control processes. Under reduced evaluative coupling, however, these configurations may re-emerge as coherent experiential states. This interpretation is consistent with contemporary accounts of memory reactivation, in which partial cues can trigger high-fidelity reinstantiation of previously encoded configurations without requiring deliberate recall or narrative reconstruction.

Notably, the permissive regime does not generate new affective content, but restores access to configurations that remain encoded yet are typically inaccessible under standard control dynamics. Accordingly, reactivation inherits the same dependency on Z and g as all other regime-level phenomena.

7 All-or-None Entry and Regime Shift Dynamics

Entry into the permissive regime is characterized by abrupt, unambiguous transitions rather than gradual changes. Probabilistic access and abrupt entry refer to distinct aspects of the same process and should not be conflated. The probability of entering the regime depends on the system's proximity to the critical threshold, which is shaped by fluctuations in evaluative load (Z). By contrast, once this threshold is crossed, the transition itself is discontinuous and experienced as abrupt. Accordingly, access is probabilistic, while entry is all-or-none. Reports frequently describe a clear onset, with little to no intermediate or uncertain states. Ambiguous or "partial" entries are rare, suggesting that access is governed by threshold dynamics rather than continuous accumulation.

Formally, this transition corresponds to the crossing of a critical boundary in evaluative load (Z), such that a reduction below threshold weakens the coupling between discrepancy and correction ($g < g^*$). Importantly, this all-or-none characterization applies to regime entry rather than to the internal dynamics of the state. Once the regime is entered, the intensity of the experience may vary continuously as a function of ongoing signal dynamics.

The phenomenological profile associated with this transition is not used here as evidence for the underlying mechanism, but as a set of constraints that competing explanations must account for. Alternative interpretations, such as immersive distraction or gradual mood elevation, generate distinct and testable predictions. The transition itself can be described as a regime shift, marked by a sudden and sustained increase in positive affect. In some cases, the resulting intensity is reported as unusually high, occasionally approaching levels that are difficult to tolerate.

Importantly, the shift is multidimensional. Alongside elevated positive affect, reports frequently include increased sensitivity to reward, enhanced emotional variability, vivid imagery, spontaneous laughter, and a general amplification of experiential salience. The state is typically described as self-sufficient and non-teleological, emerging independently of explicit goals or outcomes. Joy may present as a chest-centered pleasure, but what is noticed is a global surge in experiential intensity. A notable feature of this transition is its immediate recognizability. In some cases, the initial phase may include atypical or difficult-to-classify affective qualities, reflecting the abrupt reconfiguration of control dynamics associated with reduced evaluative load (Z) and weakened coupling (g). This pattern has been repeatedly observed in informal and exploratory settings (Morin, 2026). Taken together, these observations suggest that access to high-intensity states is not gradual, but depends on whether the system enters a distinct operational regime.

8 Joy as a Threshold-Dependent Regime Property

In the present framework, joy is not treated as a direct product of inputs, actions, or optimization processes. Instead, it is conceptualized as a regime-dependent property of the system, contingent on the level of evaluative control. This control is captured by a variable Z, which reflects the degree to which the system engages in continuous comparison, evaluation, and correction. Z indexes the strength of evaluative coupling that transforms discrepancies into goal-directed adjustments. The system is governed by a critical threshold, $Z_threshold$. When Z remains above this threshold, the system operates in an evaluation-dominated regime. When Z falls below this threshold, a qualitative shift occurs, and the system becomes permissive to a distinct regime in which high-intensity positive affect becomes accessible.

Formally:

- If $Z > Z_threshold$, the system remains in an evaluation-dominated regime
- If $Z < Z_threshold$, the system transitions into a permissive regime

Accordingly, the threshold applies to access conditions, not to the magnitude of the resulting affect. This transition reflects a discontinuous change in system dynamics rather than a gradual modulation. In the evaluation-dominated regime (Z high), processing is characterized by continuous monitoring, goal-directed behavior, performance tracking, and rapid error correction. Experience is structured by narrative integration and outcome relevance, and positive affect tends to be limited, unstable, and dependent on external contingencies. By contrast, in the permissive regime (Z low), evaluative processes fail to stabilize discrepancies. This results in a relative suspension of optimization, reduced coupling between action and instrumental outcomes, and increased sensitivity to structured input. Within this configuration, high-intensity positive affect becomes accessible as a property of the system's dynamics rather than as the result of reward accumulation.

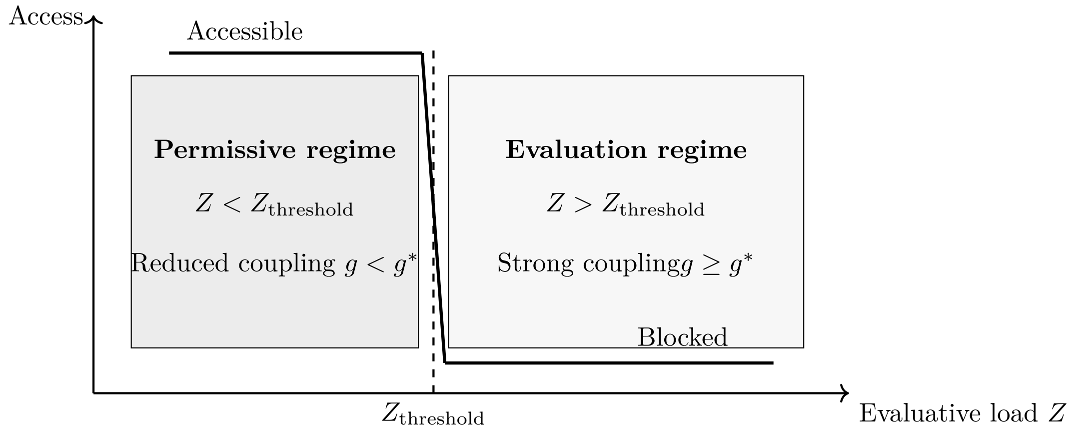

Figure 2: Threshold regime model.

9 Anti-Instrumental Access and Control Sensitivity

Within the proposed framework, attempts to intentionally use, optimize, or reproduce the state are predicted to increase evaluative load, thereby reducing the probability of regime access. More precisely, evaluation increases evaluative load (Z), which in turn restores the coupling between discrepancy and correction ($g \geq g^*$), thereby preventing or terminating regime access. Goal-directed control and monitoring processes increase Z, pushing the system above the critical threshold.

Accordingly, access to the permissive regime is disrupted by conditions that introduce explicit or implicit optimization demands, including goal setting, performance-oriented repetition, and real-time evaluation. This property distinguishes the present account from frameworks in which positive affect can be enhanced through training or reward maximization. In contrast, increasing control in this sense reduces access by increasing evaluative load (Z) and thereby strengthening coupling (g).

Interventions should therefore target the structure of control rather than affect directly, for example by reducing monitoring or limiting narrative capture. This constraint applies not only to explicit goals but also to continuous micro-level corrections applied to ongoing experience.

Within the permissive regime, external stimuli modulate the intensity of the state without generating it. Inputs that introduce structured but manageable discrepancies, such as small and frequent prediction errors without strong demands for resolution, are predicted to amplify experiential intensity. Under reduced evaluative load (Z), the weakening of coupling (g) allows these discrepancy signals to persist and contribute directly to experience. Stimuli with high variability and low requirements for causal coherence, including rapidly changing audiovisual sequences, are consistent with this profile, as they can sustain discrepancy signals without triggering large-scale corrective processes.

A key implication is that the challenge in sustaining high-intensity positive affect does not lie in identifying appropriate activities, but in preventing their instrumentalization. Once an activity becomes associated with goals, expectations, or evaluation, it increases evaluative load (Z), strengthens coupling (g), and reduces access. Moreover, the modulatory effects of stimuli depend on prior regime entry. In the absence of access, similar inputs are not expected to produce comparable effects, reinforcing the distinction between modulatory and generative factors.

From this perspective, the regime is difficult to access not because suitable stimuli are rare, but because the conditions required for access are systematically undermined by optimization.

10 Evaluative Load as a General Access Constraint

Within the proposed framework, access to the permissive regime depends on reduced evaluative coupling at the point of entry. Any process that introduces evaluation, monitoring, or goal-directed control increases evaluative load (Z), thereby strengthening the coupling between discrepancy and correction ($g \geq g^*$), and reducing or preventing regime access.

This constraint applies broadly across contexts. Explicit goal-setting, performance-oriented behavior, and attempts to intentionally induce or optimize the state all increase evaluative load. Likewise, measurement procedures such as real-time ratings, introspective prompts, or performance framing recruit self-monitoring processes that act on the same control pathway. In both cases, evaluation does not merely describe the state, but actively alters the underlying dynamics by increasing Z and restoring corrective coupling.

Accordingly, access to high-intensity positive affect is anti-instrumental in its entry conditions. Attempts to use, reproduce, or stabilize the state tend to reduce its probability of occurrence by reintroducing evaluative processes. Measurement is therefore not neutral, but functions as an intervention on control dynamics, with the potential to suppress the phenomenon under observation or produce attenuated effects.

This leads to a methodological constraint. The absence or rarity of such states in experimental settings may reflect the evaluative structure of the task rather than true inaccessibility. Empirical investigation remains possible, but valid approaches must minimize interference with regime access. In particular, indirect, post-hoc, or minimally intrusive methods are more appropriate than real-time evaluation.

The model generates several methodological predictions. First, post-hoc reporting should yield higher observed incidence than real-time assessment. Second, behavioral or indirect proxies may detect regime transitions without requiring explicit monitoring. Third, experimental designs that reduce anticipatory evaluation, for example through partial blinding to task purpose, should increase access probability by lowering Z. Finally, spontaneous reports under low evaluative conditions may serve as behavioral markers of regime entry.

More generally, these considerations imply that observation itself must be treated as a variable in experimental design. In this class of phenomena, the most direct measurement procedures may be the least valid, as they recruit the very processes that prevent the state from occurring. Converging evidence from studies comparing real-time experience sampling with retrospective reports suggests that measurement itself may act as an intervention on affective dynamics, altering both intensity and structure of reported states. This raises the possibility that certain high-intensity experiences are systematically under-detected under standard experimental conditions.

11 Monitoring, Evaluation, and Optimization

A central distinction in the present framework concerns the relationship between monitoring, evaluation, and optimization. Monitoring is defined as a control-layer process that continuously tracks discrepancies between expected and observed states. Evaluation is treated as an observable manifestation of this process, reflecting the presence of evaluative load (Z), rather than its underlying mechanism. Optimization refers to the behavioral regime that emerges when discrepancies are systematically converted into corrective or goal-directed adjustments. In this sense, optimization is not a voluntary strategy but a mode of operation induced by sustained monitoring through increased evaluative load.

Monitoring functions as a stabilizing layer that maintains the system in an evaluation-dominated regime. Mechanistically, sustained monitoring increases evaluative load (Z), which strengthens the coupling between discrepancy and correction (g), ensuring that discrepancies are rapidly transformed into goal-relevant signals. This reinforces its own activity, leading to a default mode in which experience is continuously interpreted, categorized, and adjusted.

This control structure is not directly accessible without altering it, as introspection and reporting recruit evaluative processes and thereby increase evaluative load (Z), modifying the underlying dynamics. The empirical target is therefore not monitoring itself, but its behavioral signatures, including abrupt regime entry, sensitivity to evaluative framing, disruption by goal-setting or repetition, and dissociation between access and skill acquisition.

The framework further predicts that monitoring is indirectly reinforced through its alignment with reward and meaning attribution. States associated with high evaluative load (Z) may be experienced as meaningful, competent, or task-relevant, stabilizing the optimization regime. Conversely, states with reduced evaluative load may be perceived as less legitimate, biasing the system toward reinstating monitoring. As a result, the optimization regime becomes self-stabilizing, both through control dynamics and through its compatibility with reward and meaning signals.

Accordingly, evaluation does not merely describe experience but acts by increasing evaluative load (Z), which strengthens coupling (g) and reshapes the regime in which experience unfolds, contributing to the systematic suppression of experiential modes that depend on reduced evaluative coupling.

11.1 Cultural Sources of Evaluative Load

The preceding analysis identifies monitoring, evaluation, and optimization as central drivers of evaluative load (Z). These processes are not only internally generated but are also reinforced by widely adopted cultural and theoretical frameworks. In many cases, such frameworks do not merely describe affective experience, but actively structure the way it is approached, interpreted, and engaged with.

The following table summarizes a set of common frameworks that may implicitly increase evaluative load by introducing conditional, comparative, or optimization-based relationships to experience.

Table 1: Cultural and theoretical frameworks that may implicitly increase evaluative load (Z) by inducing persistent forms of monitoring.

Framework	**Core assumption**	**Implicit effect on Z**	**Induced monitoring type**
Hedonic adaptation / habituation	Affective responses decrease with repetition	Installs expectation of decline, leading to anticipatory checking	Temporal monitoring: *is it fading?*
Emotion regulation models	Emotions should be managed and optimized	Introduces continuous top-down control over experience	Normative monitoring: *am I regulating correctly?*
Reward-based accounts	Positive affect depends on outcomes and prediction errors	Encourages outcome-seeking and comparison between expectation and result	Outcome monitoring: *did I get the reward?*
Flow (general model)	Optimal experience arises from a balance between skill and challenge	Requires ongoing alignment between internal capacity and task demands	State monitoring: *am I in flow?*
Flow (difficulty optimization variant)	Enjoyment depends on maintaining optimal challenge	Transforms experience into a calibration task requiring continuous adjustment	Calibration monitoring: *is this difficulty optimal?*

Framework	**Core assumption**	**Implicit effect on Z**	**Induced monitoring type**
Quality ranking / tier-list culture	Experiences can be ranked and optimized to identify the best option	Converts experience into a comparative selection problem with constant alternatives	Comparative monitoring: *is this the best option?*
Complexity-based happiness models	Well-being depends on optimizing multiple factors simultaneously	Turns experience into a high-dimensional constraint problem	Constraint monitoring: *are all conditions satisfied?*
Serotonin / lifestyle guarantee models	Positive mood is reliably produced by specific conditions such as sunlight or exercise	Creates expectation of predictable outcomes and mismatch detection	Causal monitoring: *is this working as expected?*
Memory-centered value theory	The value of an experience lies in the memory it will produce	Shifts focus toward future recall and narrative construction	Prospective monitoring: *will this be a good memory?*
Nostalgic idealization (*it was better before*)	Past periods were inherently richer or more authentic than the present	Devalues current experience by comparison with an idealized past, sustaining dissatisfaction and disengagement	Temporal-comparative monitoring: *was it better before?*
Measurable well-being frameworks	Well-being is stable and quantifiable	Promotes self-evaluation and conversion into metrics	Reflective monitoring: *how do I rate this?*
Meditation (instrumentalized form)	Desired states must be actively induced through technique	Reframes experience as something to produce and track	Induction monitoring: *am I entering the state?*
Meditation (long-term mastery model)	Access requires years of training and exceptional motivation	Defers access and ties it to progress and qualification	Progress monitoring: *am I advanced enough yet?*
Pharmacological exclusivity belief	Intense positive states require external substances	Externalizes access and creates dependency on triggers	Access monitoring: *do I have the required trigger?*
Elsewhere bias / context-switching belief	A better experience is always available in another context	Prevents settling by maintaining counterfactual alternatives	Counterfactual monitoring: *would it be better elsewhere?*

These frameworks converge on a common structural property: they treat affective experience as conditional, whether on outcomes, optimization, correct parameters, future value, or external triggers. This conditionalization introduces persistent forms of monitoring, including anticipatory, comparative, and reflective processes. Within the present framework, these processes correspond to increases in evaluative load (Z), which strengthen the coupling between discrepancy and correction (g), thereby reducing the probability of regime access.

Importantly, this implies that the rarity or instability of high-intensity positive states may not solely reflect intrinsic properties of the system, but also the widespread adoption of cognitive and cultural frameworks that maintain elevated evaluative load. The effectiveness of non-instrumental tasks may be modulated by pre-existing cognitive frameworks that impose evaluative or optimization-based interpretations. Even when a task is structurally designed to reduce monitoring, beliefs that frame experience as conditional, measurable, or optimizable may reintroduce evaluative load (Z), thereby reducing the probability of regime access. While the threshold model accounts for access conditions, it does not yet explain why entry remains difficult even when these conditions appear to be met.

12 Regime Collapse, Hysteresis, and the Z_{shift} Mechanism

Within the proposed framework, collapse of the permissive regime corresponds to the re-engagement of evaluative control, which restores the coupling between discrepancy detection and corrective processes and thereby terminates the regime. Empirically, collapse is often associated with meta-evaluative activity, such as monitoring the quality, duration, or meaning of the state. These processes are sufficient to alter control dynamics even in the absence of changes in external input or reward.

A key property of the system is its asymmetric stability. Entry is constrained by threshold conditions and appears fragile, whereas persistence can tolerate partial re-engagement of evaluative processes. This asymmetry is consistent with hysteresis-like dynamics, in which the authority of evaluative control depends on recent state history. Formally, entry requires suppression of evaluative gain below a critical threshold, while persistence allows partial increases without immediate collapse. Collapse, however, reinstates evaluative authority in a way that prevents rapid re-entry.

A central mechanism (Zshift) involves anticipatory monitoring that maintains elevated evaluative load and prevents re-entry. This mechanism is detailed in the following section. This mechanism suggests a potential developmental transition in which the emergence of anticipatory testing reduces spontaneous access to the permissive regime. Rather than a gradual change, this transition may reflect a discrete shift associated with the realization that the state does not reliably return when expected. Z_{shift} is best understood as a functional construct capturing a set of convergent control-related processes, including evaluative monitoring, goal-directed control, and anticipatory regulation. Among these, evaluative monitoring appears to play a central and partially controllable role, as even minimal forms of self-observation can significantly alter the probability of regime entry.

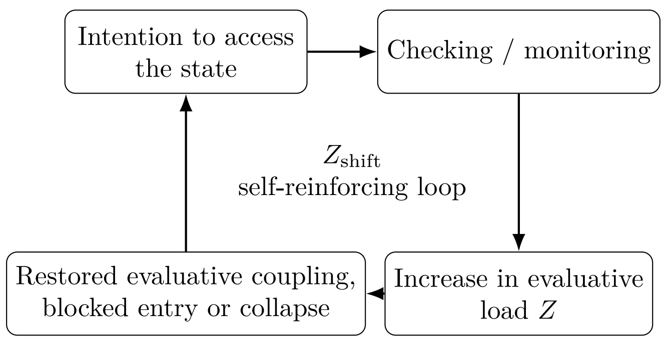

Figure 3: Schematic representation of the Z_{shift} mechanism. Attempts to access or verify the state recruit monitoring processes, which increase evaluative load (Z), restore evaluative coupling, and thereby block entry or terminate the permissive regime. This generates a self-reinforcing loop in which the intention to recover the state helps maintain the conditions that prevent its return.

Monitoring Dynamics and the Z_{shift} Mechanism

Within the proposed framework, monitoring is defined as a control-layer process that continuously tracks discrepancies between expected and observed states. Under standard conditions, this process contributes to the stabilization of behavior by converting discrepancies into evaluative signals and corrective adjustments. Sustained monitoring increases evaluative load Z, which strengthens the coupling between discrepancy and correction ($g \geq g^*$), maintaining the system in an evaluation-dominated regime.

Evaluation and optimization are treated as behavioral manifestations of this underlying monitoring process. When discrepancies are systematically interpreted in terms of performance, meaning, or goal relevance, they are rapidly transformed into corrective drives. This reinforces the monitoring process itself, producing a stable mode of operation in which experience is continuously interpreted, adjusted, and aligned with task demands.

A key implication is that monitoring is not a neutral or purely observational process. Acts such as introspection, evaluation, or attempts to control experience recruit the same control dynamics, increasing Z and modifying the regime in which processing unfolds. As a result, monitoring contributes directly to the suppression of experiential modes that depend on reduced evaluative coupling.

The $Zshift$ mechanism describes a specific dynamical extension of this process. It corresponds to a self-reinforcing loop in which attempts to detect, access, or recover the permissive regime induce anticipatory monitoring. In this loop, the intention to access the state leads to checking, and checking leads to sustained evaluative activity, maintaining elevated Z over time. This prevents the reduction of coupling required for regime entry.

Over repeated exposure, this loop may stabilize into a background process, shifting the system from intermittent monitoring to continuous anticipatory evaluation. In this configuration, monitoring is no longer driven solely by immediate discrepancies, but by expectations about future internal states. This results in a persistent elevation of Z, reducing the probability of regime access even in the absence of external evaluative demands.

This dynamic also introduces an asymmetry between entry and persistence. Entry into the permissive regime requires a sufficient reduction in evaluative load to weaken coupling ($g < g^*$), whereas persistence can tolerate partial re-engagement of monitoring without immediate collapse. However, once the Z_{shift} loop is reactivated, evaluative load increases rapidly, restoring coupling and terminating the regime. Following collapse, the system remains biased toward elevated Z, preventing rapid re-entry.

Accordingly, monitoring and $Zshift$ should be understood as components of a single control dynamic. Monitoring provides the baseline mechanism through which discrepancies are evaluated and corrected, while Z_{shift} represents a self-sustaining amplification of this process driven by anticipatory evaluation. Together, they define a control structure that stabilizes the evaluation-dominated regime and constrains access to the permissive regime.

13 Structured Prediction Error and Engagement in Cartoons

Cartoons provide a useful example of stimulus structures that may interact with the permissive regime. They are typically characterized by weak or absent resolution, minimal goal-directed structure, and low or inconsistent payoff contingencies. Rather than relying primarily on classical humor mechanisms such as setup and punchline, cartoons often exhibit dense sequences of perceptual and causal violations. Physical laws may be intermittently broken, proportions destabilized, and objects transformed without warning. These features generate sustained prediction error and perceptual instability.

From this perspective, what is commonly interpreted as humor may partly reflect an evaluative overlay imposed on more fundamental processes. An alternative account is that such stimuli operate through structured prediction error, maintaining discrepancy signals without requiring immediate resolution.

Within the present framework, such conditions are more likely to occur under reduced evaluative load (Z), where coupling between discrepancy and correction (g) is weakened. Under these conditions, prediction error can persist in an informational mode rather than being rapidly converted into corrective or evaluative processes.

Children's engagement with cartoons may therefore reflect more frequent access to such regimes. Importantly, the absence of explicit reports of strong positive affect in these contexts does not imply its absence, but is consistent with a regime in which affect is not amplified, categorized, or reported.

More generally, these observations support the claim that prediction error does not require resolution to sustain engagement. Rather, sustained engagement may depend on reduced evaluative load (Z) and the absence of enforced closure, which together allow discrepancy signals to persist without immediate correction. Their effect is contingent on reduced evaluative coupling, and is therefore not attributable to stimulus structure alone.

14 Control Reduction, Signal Coherence, and Pharmacological Dissociations

Pharmacological interventions that reduce constraint or evaluative control do not reliably produce the regime described here. Although such interventions may increase affective flexibility or reduce negative affect, they often fail to induce a coherent and stable transition into a high-intensity positive regime.

For example, ketamine, which has been associated with reduced activity in anterior cingulate and related control networks, as well as rapid antidepressant effects, does not consistently produce a sustained or structured increase in positive affect. Instead, its effects are often characterized by dissociation, perceptual instability, or diffuse affective changes.

This dissociation suggests that a reduction in evaluative control, taken in isolation, is not sufficient for regime entry. Within the present framework, this corresponds to the fact that reducing evaluative load (Z) is necessary but not sufficient, as access also depends on the integrity of discrepancy signals and their capacity to propagate coherently.

Specifically, entry requires a selective attenuation of evaluative micro-optimization, reflected in a reduction of coupling (g), within an otherwise intact perceptual and salience architecture. When control reduction is global or accompanied by signal degradation, discrepancy signals may lose coherence and fail to propagate in a structured manner.

Under such conditions, the system may exhibit reduced evaluative load (Z) and weakened coupling (g), yet fail to achieve the organized persistence of discrepancy required for high-intensity positive affect. This distinction allows the model to accommodate mechanistic evidence linking control-related structures to affective regulation, while avoiding reduction to a simple inhibition account.

In particular, the framework predicts that interventions which preserve signal coherence while reducing evaluative load (Z) and coupling (g) should be more effective in enabling regime transitions than those that broadly disrupt neural processing. More generally, this perspective suggests that the relationship between control and affect is not monotonic: reducing control can produce qualitatively different outcomes depending on whether underlying signal structure is preserved or degraded.

An additional implication concerns tolerance. Rather than reflecting only reduced reward sensitivity, tolerance effects may partly arise from anticipatory control processes, in which repeated exposure increases evaluative load (Z) through monitoring and prediction of the intervention's effects, thereby restoring coupling (g) and reducing regime accessibility.

15 The Morin Z-Reduction Task (M-ZRT): A Discriminative Protocol

The Morin Z-Reduction Task (M-ZRT) is not intended as a standard induction technique, but as a discriminative protocol designed to differentiate between two classes of models: gradual, learning-based accounts and threshold-based regime-shift accounts.

In learning-based frameworks, repeated exposure to a task is expected to produce gradual improvement, increased reproducibility, and performance-dependent effects. By contrast, the present model predicts a distinct signature: repeated failures, followed by abrupt onset, a qualitative transition, and persistence beyond the initial context. The presence of discontinuous success following unsuccessful attempts is therefore treated not as noise, but as evidence of threshold-governed access.

Importantly, the M-ZRT is constructed to avoid the stabilization of evaluative control. Attempts to directly reduce evaluation are themselves predicted to increase monitoring, creating a self- reinforcing loop that increases evaluative load (Z) and prevents regime entry. For this reason, control over monitoring must be indirect. Any protocol that becomes stable, repeatable, or optimizable is expected to increase evaluative load (Z), thereby strengthening coupling (g) and reducing its own effectiveness.

The task is designed as a brief, non-instrumental perturbation of ongoing activity. Short sequences of actions may be more effective than single manipulations, as isolated interventions may be too weak or ambiguous, whereas extended sequences risk becoming structured and subject to optimization. The aim is to disrupt motor and cognitive continuity without allowing the formation of a repeatable routine.

A key constraint of the task is the absence of explicit goals, metrics, evaluation, repetition, or expectation. These conditions are intended to minimize evaluative load (Z) and prevent the reactivation of monitoring processes, thereby weakening the coupling between discrepancy and correction (g).

As an illustrative example, participants may engage in an ongoing activity (e.g., navigating a virtual environment) while introducing brief, non-instrumental perturbations, such as generating a transient mental image without attempting to maintain it, followed by an irregular, non-rhythmic motor action. The sequence is terminated without continuation or evaluation. The task is performed sparingly, and not repeated in a structured manner.

The model further predicts specific failure modes corresponding to the reintroduction of evaluative processes. These include anticipatory commitment, premature closure, performance monitoring, meaning attribution, corrective adjustment, and expectation-driven resolution. Mechanistically, these processes increase evaluative load (Z), restore coupling ($g \geq g^*$), and are therefore expected to block or terminate regime entry. The task operates indirectly on Z, and does not attempt to manipulate affective intensity directly.

16 Perturbation Classes in the M-ZRT: Openness, Salience, and Suspension

The M-ZRT can be decomposed into a set of minimal perturbation classes that differentially target evaluative control processes. These perturbations are not defined by their content, but by their functional effect on goal closure, attentional allocation, and motor continuity. Openness refers to transient operations that weaken immediate goal closure or deliberate control. These perturbations introduce elements that are not used to advance task objectives and are not maintained or evaluated. For example, the generation of a brief, low-purpose mental image during ongoing activity may function as an openness perturbation, provided that it is neither sustained nor used instrumentally. Similarly, introducing an incomplete thought (e.g., "what if. . . ") without resolving or revisiting it disrupts closure without engaging problem-solving processes. Salience refers to the temporary amplification of a single perceptual or internal element without converting it into a goal. These perturbations increase the local prominence of a signal while avoiding integration into task-relevant structure. For instance, the spontaneous fading of a mental image may act as a salience event by briefly capturing attention through a minimal internal change.

Likewise, selecting a single sensory input (e.g., a sound) and momentarily treating it as dominant, without attaching relevance or action to it, produces a transient salience shift. Suspension refers to brief interruptions of ongoing optimization or strategic continuity. These perturbations disrupt the flow of goal-directed behavior without introducing an alternative objective. For example, a short, non-instrumental interruption of movement followed by an uncorrected change in direction can break motor continuity while avoiding compensatory adjustment or explanation. These three classes can be combined to produce composite perturbations that target evaluative coupling from multiple angles. In particular, sequences that introduce openness, followed by a salience event, and terminate with a suspension may be effective in disrupting monitoring without allowing the formation of a repeatable structure. Importantly, the target configuration is not reduced activity or disengagement, but a state in which salience is preserved while evaluative control is attenuated. The aim is therefore not to suppress processing, but to decouple salience from optimization. These perturbation classes are defined functionally and may be implemented across modalities, tasks, or environments.

16.1 Limitations and Boundary Conditions of the M-ZRT Protocol

The effectiveness of the M-ZRT protocol may be constrained by pre-existing cognitive and cultural frameworks that shape how participants interpret the task. While the protocol is designed to reduce evaluative load (Z) by removing explicit goals, structure, and optimization demands, its success depends on whether participants are able to engage with the task in a genuinely non-instrumental manner.

In many cases, widely adopted beliefs about positive affect may implicitly reintroduce evaluative processes. Frameworks that treat experience as conditional on outcomes, optimal parameters, correct execution, or measurable effects can lead participants to interpret the task as a method to be performed correctly or a procedure to be evaluated. This interpretation may occur even in the absence of explicit instructions. Such processes correspond to increases in evaluative load (Z), and may trigger the reactivation of the Z_{shift} mechanism, thereby preventing access to the target regime. It may take weeks to reverse deeply ingrained evaluative and optimization-based habits.

17 Discussion

The present framework reorients the analysis of high-intensity positive affect by shifting the explanatory focus from production mechanisms to access conditions. Rather than treating such states as outcomes of reward, learning, or optimization, the account emphasizes the role of system-level control dynamics in determining their availability. This perspective allows disparate observations, including abrupt onset, low reproducibility, and sensitivity to context, to be understood within a unified structure.

This shift has direct consequences for how competing models are evaluated. Incremental and learning-based accounts predict gradual improvement, stability through repetition, and dependence on skill acquisition. In contrast, a threshold-based framework predicts discontinuity, weak trainability, and the possibility of repeated failure preceding abrupt transitions. Empirical identification of such patterns would therefore provide a critical point of discrimination between these classes of models.

The framework also highlights a structural constraint on voluntary control. The difficulty of re-accessing high-intensity positive states does not appear to arise from insufficient stimulation or degraded reward sensitivity, but from a control dynamic in which attempts at access alter the conditions required for entry. This introduces a self-limiting structure that constrains intentional reproduction, even when underlying affective systems remain intact. Recent accounts of dopaminergic regulation and hedonic adaptation further suggest that anticipation, evaluation, and repeated exposure can attenuate affective intensity over time, potentially through mechanisms involving increased predictive control and expectation (e.g., Lembke, 2021).

A key implication concerns methodology. If access depends on reduced evaluation at entry, then standard experimental procedures, particularly those involving real-time reporting or introspective monitoring, may systematically interfere with the phenomenon under investigation. This suggests that inconsistencies in the literature may partly reflect differences in evaluative load induced by measurement protocols. Accordingly, approaches that minimize interference, including post-hoc reports, behavioral indicators, and low-intrusion designs, are better suited to capturing regime transitions.

The account also offers a reinterpretation of developmental patterns. The apparent reduction in the accessibility of high-intensity positive states across development may reflect changes in control structure rather than a loss of underlying capacity. Increased stabilization of evaluative processes would shift baseline conditions away from those that permit regime transitions, providing a parsimonious explanation for both the intensity of early experiences and their later inaccessibility. This perspective also clarifies why memory-based approaches, such as nostalgia, may fail to recover such states despite reactivating relevant content.

At the level of stimulus interaction, the framework distinguishes between factors that modulate intensity and those that determine access. Inputs characterized by structured variability or sustained prediction error may amplify ongoing states but do not appear sufficient to induce them in isolation. Their effects depend on prior access conditions, which explains the context-dependent efficacy of the same stimuli.

Several limitations remain. The construct of evaluative load requires operationalization through measurable proxies, and its relationship to specific neural systems remains to be specified. The proposed dynamics are consistent with existing accounts of salience, monitoring, and control, but their precise implementation is underspecified. In addition, the phenomenological observations motivating the framework require systematic validation under controlled conditions.

Future work should prioritize experimental designs capable of detecting discontinuous transitions while minimizing evaluative interference. The proposed behavioral protocol provides one possible approach, particularly as a means of distinguishing threshold dynamics from gradualist accounts. More broadly, progress in this domain may depend less on identifying optimal stimuli than on refining methods that preserve the conditions under which such states can occur. Taken together, these observations support the view that high-intensity positive affect may not be limited by reward availability or learning capacity, but by control-related constraints that regulate access to specific experiential regimes.

18 Implications for Development, Psychopathology, and Learning

If positive affect is partially governed by regime dynamics rather than solely by reward magnitude or valuation processes, this has broad implications across multiple domains of cognitive science

18.1 Development

The proposed framework offers a reinterpretation of developmental changes in affective experience. Early childhood is often characterized by frequent access to high-intensity positive states that appear difficult to reproduce in adulthood. Rather than attributing this difference to changes in reward sensitivity or novelty exposure alone, the present account suggests that developmental shifts in evaluative monitoring may play a central role. Specifically, the progressive installation of stable evaluative and goal-directed control during adolescence may increase the baseline coupling between discrepancy detection and correction. This would reduce the probability of entering regimes in which discrepancies remain informational. From this perspective, the apparent “loss” of intense positive affect is not due to disappearance of capacity, but to a change in access conditions driven by control architecture.

18.2 Psychopathology

The model also has implications for conditions involving altered affective regulation. Many forms of anxiety and mood disorders are associated with heightened evaluative monitoring, persistent error checking, and increased sensitivity to discrepancies. Within the present framework, such conditions can be understood as regimes in which the coupling between discrepancy and correction is excessively strong or rigid. This predicts a reduced accessibility of permissive regimes of positive affect, not because reward systems are impaired per se, but because evaluative processes prevent the stabilization of such regimes. Conversely, certain pharmacological or behavioral interventions that reduce evaluative gain may transiently increase access to these states, although potentially at the cost of reduced control in other domains. Importantly, this account distinguishes between reducing negative affect and enabling access to high-intensity positive regimes. These may involve overlapping but non-identical mechanisms, suggesting that standard anxiolytic or antidepressant approaches may not directly restore access to such states.

18.3 Reward and reinforcement learning

The framework challenges the assumption that increasing reward or positive prediction error is sufficient to produce high-intensity positive affect. If prediction error is rapidly converted into corrective signals under standard control regimes, its contribution to subjective experience may be limited. This suggests that the experiential impact of prediction error depends on its integration within the control architecture. Under reduced evaluative monitoring, prediction error may accumulate or persist in a way that enhances affective intensity. Under high evaluative gain, the same signals may be rapidly neutralized through updating and action. As a result, interventions aimed at modulating reward signals alone may have limited effects unless they also alter the coupling between discrepancy detection and correction.

18.4 Learning and exploration

Finally, the model has implications for learning and exploratory behavior. Standard accounts emphasize the role of prediction error minimization and reward maximization in guiding behavior. However, if certain regimes allow discrepancies to persist without immediate correction, this may create conditions for a different mode of exploration. In such regimes, behavior may become less constrained by immediate optimization and more driven by local salience or unfolding perceptual dynamics. This could facilitate forms of exploration that are not directly tied to goal attainment, but still contribute to learning by exposing the system to a broader range of states and transitions. This perspective suggests that alternating between regimes of tight control and regimes of reduced evaluative coupling may be beneficial for balancing exploitation and exploration. The present account isolates one such regime and provides a framework for studying its properties and transitions.

19 Predictions and Falsifiability

The present framework generates a set of empirically testable predictions that distinguish it from standard reward-based and continuous models of positive affect. These predictions diverge from standard reward-based and continuous models, which would instead predict graded increases, reliability under repetition, and compatibility with measurement.

1. Discontinuous outcome distributions

Affective outcomes following repeated exposure to similar conditions should not follow a graded or normal distribution. Instead, most trials are expected to produce no notable effect, while a minority produce sharply delineated, high-intensity episodes. This predicts a highly skewed or bimodal distribution of reported experiences.

2. Sensitivity to evaluative monitoring

Introducing explicit monitoring, evaluation, or performance pressure during an ongoing episode should reliably reduce or terminate the state. Even minimal forms of introspective checking are expected to decrease both intensity and duration.

3. Low reliability of induction despite stable conditions

Repeated attempts to reproduce the state under identical external conditions should yield inconsistent outcomes. Entry probability should remain low and variable, even when contextual parameters are held constant.

4. Incompatibility with instrumental goals

Framing an activity in terms of explicit objectives, optimization, or expected outcomes should reduce the likelihood of state onset. Conversely, conditions that minimize goal-directed structure should increase access probability.

5. Dissociation from reward magnitude

Increasing reward magnitude, frequency, or salience should not reliably increase the probability of state onset. In some cases, higher reward predictability or optimization may reduce access.

6. Post-episode fragility under retrospective evaluation

Following an episode, attempts to analyze, stabilize, or reproduce the experience should lead to a rapid increase in evaluative monitoring and a corresponding decrease in subsequent access probability.

20 Minimal Formalization of Regime Dynamics

The proposed account can be formalized as a change in the coupling between discrepancy detection and evaluative correction within a standard control architecture.

Let $\delta(t)$ denote discrepancy, or prediction error, at time t. Under standard conditions, discrepancies are rapidly transformed into corrective signals through a gain-controlled process:

$$c(t) = g \cdot \delta(t)$$

where $c(t)$ is the corrective drive and g is an evaluative gain parameter. High values of g correspond to a strong and rapid conversion of discrepancies into goal-directed adjustments.

In this formulation, Z can be interpreted as a macro-level index of evaluative load, while g denotes the local coupling parameter through which this load is expressed in the transformation of discrepancy into corrective signals. Increases in Z strengthen the coupling between discrepancy and correction, with regime transitions occurring when g crosses a critical threshold.

The central proposal is that access to the permissive regime depends on a reduction, or temporary failure, of this coupling. Specifically, when the gain falls below a critical threshold g^*, the system no longer enforces immediate correction:

$$g < g^* \;\Rightarrow\; \text{decoupling regime}$$

Importantly, evaluative monitoring feeds back directly onto the gain parameter. Acts such as evaluating, measuring, or attempting to control the state increase evaluative load, which in turn modulates the coupling strength:

$$g = g_0 + \alpha \cdot Z(t)$$

where $Z(t)$ denotes evaluative load and $\alpha > 0$. Through this relationship, monitoring does not merely describe the system's state, but actively reshapes its control dynamics.

This introduces an intrinsic instability in the permissive regime. Because evaluative monitoring increases $Z(t)$, and $Z(t)$ directly increases g, any attempt to observe, stabilize, or reproduce the state feeds back into the control architecture that prevents its own existence. As g increases, discrepancy signals are more rapidly transformed into corrective drives, reducing their persistence and thereby suppressing their contribution to subjective intensity.

Formally, when $g \geq g^*$, the system re-enters the evaluation-dominated regime, in which discrepancies are immediately corrected and the permissive dynamics cannot be sustained. The regime is therefore self-terminating under monitoring: the very processes required for deliberate control or measurement are sufficient to restore coupling and collapse the state.

This feedback loop also provides a mechanistic account of the Z_{shift} phenomenon. Attempts to detect or recover the state generate anticipatory monitoring, which maintains elevated evaluative load over time. This sustained increase in $Z(t)$ stabilizes g above threshold, preventing re-entry even in the absence of external constraints.

Once the regime is entered ($g < g^*$), discrepancies are no longer eliminated but persist as informational signals. Their temporal dynamics are consequently altered. Rather than being rapidly damped, discrepancy signals accumulate or remain active over extended intervals:

$$\delta_{\text{eff}}(t) = \int_0^t \delta(\tau)\, w(\tau)\, d\tau$$

where $w(\tau)$ is a persistence weighting function reflecting reduced corrective attenuation.

Subjective intensity $I(t)$ is then hypothesized to depend on this effective discrepancy signal:

$$I(t) \propto \delta_{\text{eff}}(t)$$

Under high evaluative gain, that is, $g \geq g^*$, $\delta_{\text{eff}}(t)$ remains low due to rapid correction, resulting in limited experiential impact. By contrast, under low gain, that is, $g < g^*$, persistent discrepancies contribute directly to subjective intensity.

This feedback structure captures three core properties of the proposed regime. First, threshold dynamics: small variations in evaluative gain around g^* produce discontinuous changes in system behavior. Second, persistence of discrepancy: prediction error is no longer transient but becomes a sustained signal contributing directly to experience. Third, a self-terminating structure: the very processes that would enable intentional control also reintroduce coupling, rendering the regime inherently unstable under monitoring.

This minimal formulation does not specify a neural implementation, but it is consistent with architectures in which discrepancy signals, such as prediction error, and control signals, such as evaluative or salience-related processes, are separable yet dynamically coupled.

21 Conclusion

The present paper proposes that high-intensity positive affect is not adequately captured by models that treat it as a direct outcome of reward, learning, or goal-directed optimization. Instead, joy is conceptualized as a regime-dependent property of the system, emerging when evaluative coupling falls below a critical threshold and discrepancies remain informational rather than being rapidly converted into corrective demands. This framework accounts for a set of observations that are difficult to reconcile with existing approaches, including abrupt and all-or-none transitions, sensitivity to evaluative context, limited trainability, and the paradoxical effects of measurement. It also provides a unifying interpretation of phenomena spanning development, memory reactivation, perceptual engagement, and pharmacological dissociations, by linking them to a common control parameter governing the relationship between discrepancy detection and evaluative correction.

By introducing a regime-based perspective, the model shifts the focus from the production of positive affect to the conditions that enable its accessibility. In this view, the central question is not how to increase reward or optimize behavior, but how to modulate the structure of control such that discrepancies can persist and propagate within the system. Importantly, the framework is designed to be empirically vulnerable. It generates specific predictions that distinguish it from gradualist and reward-based accounts, particularly regarding discontinuity of onset, disruption by evaluation, and the absence of reliable training effects.

These predictions provide a basis for direct experimental tests, as well as for methodological revisions that minimize interference with the phenomena under study. More broadly, the proposal highlights a potential blind spot in affect research: phenomena that depend on reduced evaluative control may be systematically underrepresented due to the very methods used to investigate them. Addressing this limitation may require not only new models, but new experimental strategies that treat observation itself as a variable rather than a neutral process. Taken together, these considerations suggest that joy, in at least one of its forms, is not best understood as an output to be maximized, but as a regime to which access must be permitted.

The author declares that there are no conflicts of interest.

22 references

Barto, A. G., & Sutton, R. S. (2018). Reinforcement learning: An introduction (2nd ed.). MIT Press.

Beeman, M., & Kounios, J. (2009). The Aha! moment: The cognitive neuroscience of insight. Current Directions in Psychological Science, 18(4), 210–216.

Beilock, S. L. (2010). Choke: What the secrets of the brain reveal about getting it right when you have to. Free Press.

Berridge, K. C. (2003). Pleasures of the brain. Brain and Cognition, 52(1), 106–128.

Berridge, K. C., & Kringelbach, M. L. (2015). Pleasure systems in the brain. Neuron, 86(3), 646–664.

Berridge, K. C. (2018). Evolving concepts of emotion and motivation. Frontiers in Psychology, 9, 1647.

Campbell, D. T. (1957). Factors relevant to the validity of experiments in social settings. Psychological Bulletin, 54(4), 297–312.

Clark, A. (2013). Whatever next? Predictive brains, situated agents, and the future of cognitive science. Behavioral and Brain Sciences, 36(3), 181–204.

Clark, A. (2019). Surfing uncertainty: Prediction, action, and the embodied mind. Oxford University Press.

Csikszentmihalyi, M. (1990). Flow: The psychology of optimal experience. Harper Row.

Diener, E. (2000). Subjective well-being: The science of happiness and a proposal for a national index. American Psychologist, 55(1), 34–43.

Epley, N., & Whitchurch, E. (2008). Mirror, mirror on the wall: Enhancement in self-recognition. Personality and Social Psychology Bulletin, 34(9), 1159–1170.

Friston, K. (2010). The free-energy principle: A unified brain theory? Nature Reviews Neuroscience, 11(2), 127–138.

Gopnik, A. (2009). The philosophical baby: What children's minds tell us about truth, love, and the meaning of life. Farrar, Straus and Giroux.

Harris, P. L. (2000). The work of the imagination. Blackwell.

Kahneman, D. (1999). Objective happiness. In D. Kahneman, E. Diener, & N. Schwarz (Eds.), Well-being: The foundations of hedonic psychology (pp. 3–25). Russell Sage Foundation.

Keltner, D., & Haidt, J. (2003). Approaching awe, a moral, spiritual, and aesthetic emotion. Cognition and Emotion, 17(2), 297–314.

Kelley, H. H. (1967). Attribution theory in social psychology. In D. Levine (Ed.), Nebraska symposium on motivation (Vol. 15, pp. 192–238). University of Nebraska Press.

Lembke, A. (2021). Dopamine nation: Finding balance in the age of indulgence. Dutton.

Maslow, A. H. (1964). Religions, values, and peak-experiences. Ohio State University Press.

Morin, F. (2026). Joy is fragile under evaluation: A threshold model of entry. Figshare.

https://doi.org/10.6084/m9.figshare.31385326.v3

Rutledge, R. B., Skandali, N., Dayan, P., & Dolan, R. J. (2014). A computational and neural model of momentary subjective well-being. Proceedings of the National Academy of Sciences, 111(33), 12252–12257.

Scherer, K. R. (2009). The dynamic architecture of emotion: Evidence for the component process model. Cognition and Emotion, 23(7), 1307–1351.

Schultz, W. (2016). Dopamine reward prediction error coding. Dialogues in Clinical Neuroscience, 18(1), 23–32.

Schultz, W. (2021). Dopamine reward prediction error coding: A two-component response. Nature Reviews Neuroscience, 22, 483–495.

Schooler, J. W. (2002). Re-representing consciousness: Dissociations between experience and meta-consciousness. Trends in Cognitive Sciences, 6(8), 339–344.

Schooler, J. W. (2011). Introspecting in the spirit of William James: Comment on Fox et al. Psychological Bulletin, 137(2), 345–350.

Wilson, T. D., Schooler, J. W. (1991). Thinking too much: Introspection can reduce the quality of preferences and decisions. Journal of Personality and Social Psychology, 60(2), 181–192.