

Transient Disengagement of Evaluative Monitoring as a Precondition for Joy

A Threshold Model of Joy Entry

Florian Morin

Independent Researcher

Canonical version: <https://florianmorin.com/papers/Threshold-Constraints.html>

Abstract

Joy may depend less on reward intensity per se than on a transient reduction in evaluative monitoring at the moment positive salience emerges. I propose that a control mode associated with conflict monitoring, comparison, cost evaluation, and self-checking can block entry into a distinct affective regime even when downstream valuation remains intact. In this account, positive affect emerges when valuation is briefly permitted to unfold without immediate arbitration by monitoring processes, and it collapses when those processes re-enter the loop. This framework generates three testable predictions. First, states described as hedonic tolerance may often reflect faster re-engagement of evaluative control rather than weaker reward responses. Second, low-stakes, non-rhythmic, continuously interactive environments may prolong affect by delaying monitoring stabilization. Third, attempts to measure or optimize the state during induction may suppress the phenomenon itself.

References to the anterior cingulate cortex are offered as a mechanistically plausible illustration, not as a necessary anatomical commitment. The model is intended as a compact behavioral and experimental framework for studying access constraints on positive affect.

In this paper, *ease* denotes the specific regime underlying what is described phenomenologically as *joy*. It is made possible when evaluative pressure becomes temporarily irrelevant, with a central distinction between access, stabilization, and propagation of the episode. Because the process is threshold-dependent, a small number of failed attempts is expected before transition into a high-salience regime.

Opening

This paper advances a simple claim: access to intense positive affect is constrained by evaluative monitoring. More specifically, a control mode associated with conflict monitoring, comparison, cost evaluation, and self-checking appears to block entry into a distinct affective regime. The central proposal is not that reward intensity is absent, but that it cannot fully express itself while valuation is continuously supervised.

The anterior cingulate cortex (ACC) provides a useful mechanistic illustration because it is repeatedly implicated in error detection, conflict monitoring, effort evaluation, and persistence. However, the argument does not depend on a strict anatomical identification. The ACC is treated here as one possible instantiation of a broader evaluative-control principle.

Across repeated observations, five claims appear central.

1. Positive Affect

High positive affect appears to coincide with a brief window in which evaluative monitoring is functionally disengaged, plausibly involving reduced ACC-linked control. During this window, valuation appears to be handled downstream (striatal / limbic, with vmPFC involvement) without being rapidly contextualized or corrected.

2. Positive affect collapse

Positive affect fades not primarily due to hedonic saturation or neurotransmitter depletion, but when the ACC re-enters the loop and begins:

- evaluating the state (“this is good”),
- comparing it to expectations or alternatives,
- predicting duration, consequences, or meaning.

This re-entry appears sufficient to flatten the experience even when sensory input and dopaminergic tone remain unchanged.

3. Persistence conditions

Positive affect persists longer when the ACC is prevented from stabilizing into a rhythmic evaluative mode. Non-rhythmic motor variability, continuous but low-stakes task engagement, or perceptual flow (especially visual motion) appear to delay full ACC re-engagement without inducing distraction or cognitive load.

Notably, this state subjectively resembles a reactivation of what could be described as a transient re-entry into a childhood-like valuation regime: a mode in which stimuli are treated as intrinsically salient without reference to optimization, comparison, or narrative coherence. Importantly, this is not regression in content or cognition, but a reversion to a simpler valuation policy.

4. Tolerance reinterpretation

What is commonly described as “tolerance” may, in some cases, reflect accelerated ACC re-engagement rather than diminished reward system responsiveness. With repetition, the ACC appears to learn to anticipate the state and pre-emptively contextualize it, preventing full emergence rather than reducing magnitude.

This could explain why some individuals report that positive affect feels “blocked” rather than weaker, and why efforts to increase intensity often fail.

5. **Dissociation between feeling and value**

A dissociation emerges in which certain stimuli retain high *value* judgments despite producing weak immediate affect. This suggests that vmPFC-based valuation can remain intact while ACC-mediated salience gating suppresses felt Positive affect . In this framing, Positive affect requires not just value, but permission for value to be *felt* without supervision.

Taken together, this suggests that positive affect is not simply a function of reward intensity, but a fragile dynamical state requiring:

- intact downstream valuation,
- temporary suspension of ACC oversight,
- and brief access to a low-monitoring, early-style valuation circuit.

This framing may help explain:

- why novelty and unpredictability matter more than intensity,
- why some non-narrative or flow-based stimuli retain emotional efficacy with age,
- why pharmacological enhancement often fails when paired with excessive self-monitoring.

It is a compact, testable model, approachable via perturbation, task design emphasizing monitoring suppression, or fine-grained temporal analysis of affective onset. Early low-monitoring play may provide an intuitive analogue: low time-checking, low metric-tracking, low explicit expectation, and permissive motor variability.

Failure Modes of Openness

Anticipatory commitment occurs when the participant decides too early. This reduces uncertainty, but also reintroduces evaluative control. Examples include predicting when the effect should appear, choosing too early between alternatives, or silently checking whether the protocol is working.

Premature closure: Instead of letting the task remain active and unresolved, the participant starts looking for an endpoint, a confirmation, or a sense that something has been completed. This introduces a goal of resolution. In this account, that shift may block the emergence of the target state.

Monitoring or performance-checking is the reflex to ask, “Am I doing it right?”. Rather than remaining inside the task, the participant begins to supervise it from outside.

Meaning-generation refers to the reflexive tendency to assign a reason, meaning, or purpose to an action too quickly. Questions such as “Why did I do that?” transform the task from an enacted process into an interpreted object.

Correction or micro-adjustment refers to the reflex to “fix it just a little” after an action is completed. Instead of letting the action stand, the participant slightly revises, refines, or aligns it.

Salience veto, or down-regulation of specialness, refers to the reflexive tendency to neutralize an emerging sense of distinctiveness as soon as it appears. This may occur through explanation, comparison, or deliberate attempts at reproduction.

Expectancy resolution refers to the reflexive tendency to assume that once an expectancy state has been created, it must be resolved by an identifiable outcome. Once the participant begins searching for the outcome that will justify the anticipation, evaluative monitoring is likely to re-enter and alter the process.

FPS (first-person shooter) protocol

Fast-paced first-person shooters such as Unreal Tournament 99 or Quake may be especially effective because they combine continuous attentional capture with frequent prediction errors and a partial suspension of explicit optimization.

Very low-dose coffee can be used as a threshold helper. The dose should remain minimal, so that it does not become salient enough to trigger monitoring or explicit interpretation as a regulatory aid. It is preferable not to use it every time, as repeated use may be incorporated into the monitor.

M-NUT provides a set of exercises, including warm-up, outside-the-task, and in-task exercises. The following protocol therefore presents one possible selection.

M-NUT can be implemented using any low-stakes environment satisfying constraints C1-C5. C1: no performance feedback, C2: no optimization goal, C3: low salience, C4: distraction-compatible, C5: monitoring stop rule. Any explicit performance-check or feeling-check operation terminates the session (or marks it as invalid).

Participants are encouraged to minimize high-load obligations and socially evaluative contexts on test days. Intrusive imagined social scenarios, such as arguments or anticipated confrontations, should also be reduced where possible. More broadly, avoiding news consumption, scrolling, comment reading, and exposure to social-media metrics may help limit re-engagement of evaluative monitoring, as may reducing repeated time-checking and unnecessary muscular tension, including jaw tension.

Repeated exposure to the protocol may decrease, rather than increase, the probability of entry. For this reason, sessions are spaced by approximately 24 hours, not as a training schedule, but to reduce procedural carryover and allow evaluative monitoring to settle between attempts.

The first successful transition may occur only after several null sessions. Within a threshold framework, such delay is not interpreted as gradual training progress, but as evidence of an entry barrier. Repeated failure followed by a sudden discontinuity is therefore treated as expected behavior rather than as an anomaly.

The task : Unreal tournament 1999 with HUD off.

1. Low instrumental stakes
The action must not carry meaningful cost, reward, or consequence for the participant. Once success matters, evaluative control is likely to re-engage.
2. Brief suspension without explicit deliberation
The task may introduce a short interruption, hesitation, or deviation, but not in a way that recruits sustained comparison, planning, or conscious problem-solving.
3. No performance criterion
There must be no internal standard of doing the task “correctly.” If the participant begins to ask whether the procedure is being executed properly, the relevant condition is lost.
4. No immediate self-evaluation
The task must not be followed by checking, interpretation, or retrospective appraisal. Post hoc thoughts such as “was that right?” or “did it work?” count as invalidating resumptions of monitoring.
5. Minimal predictability without explicit surprise-seeking
The implementation may contain mild arbitrariness or non-optimization, but it should not become a deliberate search for novelty, intensity, or effect.
6. Non-rhythmic and non-ritualized execution
If repeated, the action must not stabilize into a rhythm, routine, or familiar induction sequence. Ritualization converts perturbation into a monitored technique.
7. Single-step permissiveness
The participant should be able to perform the action and move on without correction, repetition, or refinement. Re-doing, improving, or adjusting the act reintroduces evaluative control.
8. Non-dependence on task-specific content
No specific game, object, movement, or setting is theoretically essential. Any implementation is only an example if it satisfies the same functional constraints.
9. Absence of explicit expectation of effect
The participant should not treat the action as a guaranteed trigger. Strong expectation, countdown logic, or “the effect should occur now” framing

recruits anticipatory monitoring.

10. Invalidity under monitoring capture

Any implementation becomes theoretically irrelevant once it is used in a tracked, measured, optimized, or repeatedly tested manner. At that point, the procedure no longer relaxes the bottleneck, it becomes part of it.

Accordingly, the framework does not propose a fixed set of techniques. It proposes a class of admissible manipulations defined by functional constraints. Surface-level examples are expendable; the underlying requirement is a transient reduction in evaluative monitoring without replacement by a new monitored procedure.

Relation to ACC DBS findings

Recent work suggests that optimized deep brain stimulation configured to inhibit the anterior cingulate cortex (ACC) can produce antidepressant-like effects in mice, while ACC excitation can induce anhedonia (Yuan et al., 2025). Directionally, these results support the idea that a high “ACC control mode” can be sufficient to compress access to positive affect, and that lowering this mode can reopen a richer affective regime.

Ease does not propose an invasive intervention, but a behavioral framework that targets the same axis. The central variable, Z (evaluation and monitoring load, including micro-optimization), describes how self-directed control loops recruit ACC-like dynamics and gate entry into certain hedonic regimes. Ease protocols (non-use tasks, anti-preparation, last-second choices, removal of explicit success criteria) aim to reduce Z, and therefore reduce functional engagement of ACC-like control, without a performance objective.

What Ease adds beyond a simple “ACC inhibition” reading is an entry theory. It explains why attempting to reproduce, stabilize, measure, or optimize the state can be enough to make it disappear, and why standard experimental paradigms may miss the phenomenon by reintroducing Z. In this view, DBS findings provide mechanistically compatible neighbor evidence, while Ease provides behavioral levers and human-testable predictions.

Yuan, Z., Yang, H., Wang, P., Hou, X., Xu, K., Zhou, Y., Dai, R., Gao, Y., Gao, X., Guo, Q., Li, Y., Zhang, J., Mao, Z., & Luo, M. (2025). Optimized deep brain stimulation for anterior cingulate cortex inhibition produces antidepressant-like effects in mice. *Neuron*, 113(20), 3363-3373.e4. <https://doi.org/10.1016/j.neuron.2025.07.018>

Appendix 1: Illustrative Behavioral Implementations

The examples below are illustrative only. No specific game, movement, object, or environment is theoretically required. What matters is not the surface form of the action, but whether it transiently reduces evaluative monitoring without becoming a new object of optimization, checking, or expectation.

Any admissible implementation should satisfy the following functional constraints: it should be low-stakes, brief, non-instrumental, and non-rhythmic; it should not recruit explicit comparison, problem-solving, or performance-checking; it should not be followed by retrospective evaluation; and it should not be repeated often enough to become ritualized or monitorized. Once the participant begins tracking whether the action is being performed correctly or whether it is “working,” the relevant condition is lost.

Example 1. Brief **micro-movement** perturbation

A small, non-rhythmic movement may be introduced during ongoing activity, for example a slight shoulder shift, hand repositioning, or posture adjustment, provided that it is not performed as a technique and is not repeated in a patterned way. The purpose is not to create stimulation as such, but to introduce a low-stakes deviation without subsequent monitoring or correction.

Example 2. Brief **non-optimizing hesitation** at a low-stakes fork

At a minor choice point, the participant may allow a brief pause before continuing, without entering explicit comparison or trying to identify the better option. The pause is not a deliberation period. It is admissible only if action resumes without post hoc checking, self-commentary, or corrective replay.


Example 3. Low-stakes **arbitrary selection**

Where a trivial choice is available, one option may be selected arbitrarily rather than through optimization. The relevant feature is not randomness itself, but the temporary suspension of value-weighting, justification, and internal performance criteria. If the selection becomes an object of self-observation or interpretation, the example no longer satisfies the intended constraints.

Appendix 2

Historical note: some early exploratory sessions included low-dose alcohol as a mid-task condition. This was not retained in the current version of the task, because the present framework treats behavioral structure, rather than pharmacological facilitation, as the primary variable of interest.

Appendix 3



Canonical structure of the M-NUT effect (Morin, 2026). The monitoring and evaluation loop blocks entry into the ease regime by increasing performance checking and micro-optimization. The Morin Non-Use Task (M-NUT) disrupts monitoring long enough to permit a threshold-like transition. The predicted signature is repeated clean failures followed by an abrupt discontinuity, cross-context generalization, and persistence beyond acute stimulation.

In this framework, joy is not merely a complex system property, but a thresholded dynamic regime whose access depends on monitoring constraints and produces a discontinuous, falsifiable transition signature.